

Lab 1

Hubert B. Majewski

11:59PM February 18, 2021

You should have RStudio installed to edit this file. You will write code in places marked “TO-DO” to complete the problems. Some of this will be a pure programming assignment. The tools for the solutions to these problems can be found in the class practice lectures. I want you to use the methods I taught you, not for you to google and come up with whatever works. You won’t learn that way.

To “hand in” the homework, you should compile or publish this file into a PDF that includes output of your code. Once it’s done, push by the deadline to your repository in a directory called “labs”.

- Print out the numerical constant pi with ten digits after the decimal point using the internal constant pi.

```
options(digits = 11)
pi
```

```
## [1] 3.1415926536
```

- Sum up the first 103 terms of the series $1 + 1/2 + 1/4 + 1/8 + \dots$

```
sum(1/(2^(0:102)))
```

```
## [1] 2
```

- Find the product of the first 37 terms in the sequence $1/3, 1/6, 1/9 \dots$

```
prod(1/(seq(from=3, by=3, length.out=37)))
```

```
## [1] 1.613528728e-61
```

- Find the product of the first 387 terms of $1 * 1/2 * 1/4 * 1/8 * \dots$

```
prod(1/(2^(0:386)))
```

```
## [1] 0
```

Is this answer *exactly* correct?

Experienced a numerical underflow. Eventually the computer multiplied a small number by a even smaller number

- Figure out a means to express the answer more exactly. Not compute exactly, but express more exactly.

```
sum(log(base=2, 1/2^(0:386)))
```

```
## [1] -74691
```

- Create the sequence $x = [\text{Inf}, 20, 18, \dots, -20]$.

```
x <- c(Inf, seq(from=20, by = -2, to=-20))
```

Create the sequence $x = [\log_3(\text{Inf}), \log_3(100), \log_3(98), \dots, \log_3(-20)]$.

```
x <- c(Inf, seq(from=100, by = -2, to=-20))  
x <- log(base=3, x)
```

```
## Warning: NaNs produced
```

Comment on the appropriateness of the non-numeric values.

Non numeric values occur when a function does not define an output for a given input. (i.e $\log(-1) = \text{NA}$)

- Create a vector of booleans where the entry is true if $x[i]$ is positive and finite.

```
y <- x[is.finite(x) & !is.nan(x) & x > 0]
```

- Locate the indices of the non-real numbers in this vector. Hint: use the `which` function. Don't hesitate to use the documentation via `?which`.

```
which(y == FALSE)
```

```
## integer(0)
```

- Locate the indices of the infinite quantities in this vector.

```
which(is.infinite(x))
```

```
## [1] 1 52
```

- Locate the indices of the min and max in this vector. Hint: use the `which.min` and `which.max` functions.

```
which.min(x)
```

```
## [1] 52
```

```
which.max(x)
```

```
## [1] 1
```

```
which.min(y)
```

```
## [1] 50
```

```
which.max(y)
```

```
## [1] 1
```

- Count the number of unique values in `x`.

```
length(unique(x))
```

```
## [1] 53
```

```
length(unique(y))
```

```
## [1] 50
```

- Cast `x` to a factor. Do the number of levels make sense?

```
as.factor(x)
```

```
## [1] Inf 4.19180654857877 4.1734172518943 4.15464876785729
## [5] 4.13548512895119 4.11590933734319 4.09590327428938 4.07544759935851
## [9] 4.05452163806914 4.03310325630434 4.01116871959141 3.98869253500376
## [13] 3.96564727304425 3.94200336638929 3.91772888178973 3.89278926071437
## [17] 3.86714702345081 3.84076143030548 3.81358809221559 3.78557852142874
## [21] 3.75667961082847 3.72683302786084 3.69597450568212 3.66403300987579
## [25] 3.63092975357146 3.59657702661571 3.56087679500731 3.52371901428583
## [29] 3.48497958377173 3.44451784578705 3.40217350273288 3.3577627814323
## [33] 3.31107361281783 3.26185950714291 3.20983167673402 3.15464876785729
## [37] 3.09590327428938 3.03310325630434 2.96564727304425 2.89278926071437
## [41] 2.8135880922156 2.72683302786084 2.63092975357146 2.52371901428583
## [45] 2.40217350273288 2.26185950714291 2.09590327428938 1.89278926071437
## [49] 1.63092975357146 1.26185950714291 0.630929753571457 -Inf
## [53] NaN NaN NaN NaN
## [57] NaN NaN NaN NaN
## [61] NaN NaN
## 53 Levels: -Inf 0.630929753571457 1.26185950714291 ... NaN
```

- Cast `x` to integers. What do we learn about R's infinity representation in the integer data type?

```
as.integer(x[is.finite(x) & !is.nan(x)])
```

```
## [1] 4 4 4 4 4 4 4 4 4 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 2
## [39] 2 2 2 2 2 2 2 2 1 1 1 0
```

- Use `x` to create a new vector `y` containing only the real numbers in `x`.

```
y <- x[is.finite(x) & !is.nan(x)]
```

- Use the left rectangle method to numerically integrate x^2 from 0 to 1 with rectangle width size $1e-6$.

```
sum(seq(from = 0, to = 1-1e-6, by = 1e-6)^2) * 1e-6
```

```
## [1] 0.33333283333
```

- Calculate the average of 100 realizations of standard Bernoullis in one line using the `sample` function.

```
mean(sample(c(0,1), size = 1, replace=TRUE))
```

```
## [1] 0
```

- Calculate the average of 500 realizations of Bernoullis with $p = 0.9$ in one line using the `sample` and `mean` functions.

```
mean(sample(c(0, 1), size=500, replace=TRUE, prob= c(0.9,0.1)))
```

```
## [1] 0.09
```

- Calculate the average of 1000 realizations of Bernoullis with $p = 0.9$ in one line using `rbinom`.

```
mean(rbinom(1000, 1, c(0.9, 0.1)))
```

```
## [1] 0.504
```

- In class we considered a variable `x_3` which measured “criminality”. We imagined $L = 4$ levels “none”, “infraction”, “misdemeanor” and “felony”. Create a variable `x_3` here with 100 random elements (equally probable). Create it as a nominal (i.e. unordered) factor.

```
x_3 <- as.factor(sample(c("none", "infraction", "misdemeanor", "felony"), 100, replace=TRUE))
```

- Use `x_3` to create `x_3_bin`, a binary feature where 0 is no crime and 1 is any crime.

```
x_3_bin <- x_3 != "none"
```

- Use `x_3` to create `x_3_ord`, an ordered factor variable. Ensure the proper ordinal ordering.

```
x_3_ord <- as.factor(x_3)
lv <- c("none", "infraction", "misdemeanor", "felony")
x_3_ord <- factor(x_3, levels=lv, ordered=TRUE)
x_3_ord
```

```
## [1] misdemeanor infraction none felony none infraction
## [7] misdemeanor none misdemeanor none none infraction
## [13] felony felony misdemeanor infraction felony infraction
## [19] felony infraction felony misdemeanor none misdemeanor
## [25] felony felony misdemeanor infraction infraction misdemeanor
## [31] misdemeanor felony infraction felony none felony
## [37] none infraction infraction none infraction none
## [43] none none none infraction felony misdemeanor
## [49] infraction infraction misdemeanor infraction none none
## [55] felony misdemeanor none infraction infraction felony
## [61] none none infraction felony misdemeanor misdemeanor
## [67] felony felony felony none misdemeanor misdemeanor
## [73] misdemeanor felony misdemeanor misdemeanor misdemeanor none
## [79] felony none none misdemeanor misdemeanor none
## [85] felony none none none misdemeanor none
## [91] infraction none felony misdemeanor felony felony
## [97] misdemeanor misdemeanor infraction misdemeanor
## Levels: none < infraction < misdemeanor < felony
```

- Convert this variable into three binary variables without any information loss and put them into a data matrix.

```
a <- as.numeric(x_3_ord == "infraction")
b <- as.numeric(x_3_ord == "misdemeanor")
c <- as.numeric(x_3_ord == "felony")
z = c(a, b ,c)
z <- matrix(z, nrow=100, ncol=3)

head(z)
```

```
##      [,1] [,2] [,3]
## [1,]  0   1   0
## [2,]  1   0   0
## [3,]  0   0   0
## [4,]  0   0   1
## [5,]  0   0   0
## [6,]  1   0   0
```

- What should the sum of each row be (in English)?

Take a row and add each element of column and append it to a matrix global variable.

Verify that.

```
sums1 = matrix(NA, nrow = 1, ncol=nrow(z))

for (i in 1:nrow(z)) {

  sum1 <- 0

  #For each column index do sum
  for (j in 1:ncol(z)) {
```

```

    sum1 <- sum1 + z[i, j]
  }

  sums1[1,i] <- sum1
}

#sums1

#or

sums1 <- rowSums(z)
sums1

```

```

##    [1] 1 1 0 1 0 1 1 0 1 0 0 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 1 0
##   [38] 1 1 0 1 0 0 0 0 1 1 1 1 1 1 1 1 0 0 1 1 0 1 1 1 0 0 1 1 1 1 1 1 0 1 1 1 1
##   [75] 1 1 1 0 1 0 0 1 1 0 1 0 0 0 1 0 1 0 1 1 1 1 1 1 1 1 1 1

```

- How should the column sum look (in English)?

Take a column and add each element of row and append it to a matrix global variable.

Verify that.

```

sums1 = matrix(NA, nrow = 1, ncol=3)

for (j in 1:3) {

  sum1 <- 0

  #For each row index do sum
  for (i in 1:nrow(z)) {

    sum1 <- sum1 + z[i, j]

  }

  sums1[1,j] <- sum1
}

#sums1

#or

colSums(z)

```

```

## [1] 21 27 24

```

- Generate a matrix with 100 rows where the first column is realization from a normal with mean 17 and variance 38, the second column is uniform between -10 and 10, the third column is poisson with mean 6, the fourth column in exponential with lambda of 9, the fifth column is binomial with $n = 20$ and $p = 0.12$ and the sixth column is a binary variable with exactly 24% 1's dispersed randomly. Name the rows the entries of the `fake_first_names` vector.

```
fake_first_names = c(
  "Sophia", "Emma", "Olivia", "Ava", "Mia", "Isabella", "Riley",
  "Aria", "Zoe", "Charlotte", "Lily", "Layla", "Amelia", "Emily",
  "Madelyn", "Aubrey", "Adalyn", "Madison", "Chloe", "Harper",
  "Abigail", "Aaliyah", "Avery", "Evelyn", "Kaylee", "Ella", "Ellie",
  "Scarlett", "Arianna", "Hailey", "Nora", "Addison", "Brooklyn",
  "Hannah", "Mila", "Leah", "Elizabeth", "Sarah", "Eliana", "Mackenzie",
  "Peyton", "Maria", "Grace", "Adeline", "Elena", "Anna", "Victoria",
  "Camilla", "Lillian", "Natalie", "Jackson", "Aiden", "Lucas",
  "Liam", "Noah", "Ethan", "Mason", "Caden", "Oliver", "Elijah",
  "Grayson", "Jacob", "Michael", "Benjamin", "Carter", "James",
  "Jayden", "Logan", "Alexander", "Caleb", "Ryan", "Luke", "Daniel",
  "Jack", "William", "Owen", "Gabriel", "Matthew", "Connor", "Jayce",
  "Isaac", "Sebastian", "Henry", "Muhammad", "Cameron", "Wyatt",
  "Dylan", "Nathan", "Nicholas", "Julian", "Eli", "Levi", "Isaiah",
  "Landon", "David", "Christian", "Andrew", "Brayden", "John",
  "Lincoln"
)

a <- rnorm(100, mean=17, sqrt(38))
b <- runif(100, -10, 10)
c <- rpois(100, 6)
d <- rexp(100, 9)
e <- rbinom(100, 20, prob=.12)
f <- sample(c(0,1), 100, replace=TRUE, prob = c(1-.24, .24))

z <- matrix(c(a,b,c,d,e,f), nrow=100, ncol=6, byrow=FALSE)

rownames(z) <- fake_first_names

head(z)
```

```
##           [,1]      [,2] [,3]      [,4] [,5] [,6]
## Sophia    5.5254997383 -2.12655299809    5 0.226709263415    1    0
## Emma     13.2149314377  8.02768824156    5 0.044510270728    3    0
## Olivia    27.6173504444  0.67036652938    6 0.029761390140    2    1
## Ava        9.2607754556 -9.45537828840    3 0.102753815354    3    1
## Mia       12.5159214907 -3.95529460628    3 0.008654872059    1    0
## Isabella   5.3157299149  4.48686417658    7 0.177954834566    1    0
```

- Create a data frame of the same data as above except make the binary variable a factor “DOMESTIC” vs “FOREIGN” for 0 and 1 respectively. Use RStudio’s `View` function to ensure this worked as desired.

```
z1 <- data.frame(z)

for (i in 1 : nrow(z1)) {
```

```

for(j in 1: ncol(z1)) {
  if (j == 6) {
    if (z1[i,j] == 1) z1[i,j] <- "FOREIGN"
    else z1[i,j] <- "DOMESTIC"
  }
}
}

View(z1)

```

- Print out a table of the binary variable. Then print out the proportions of “DOMESTIC” vs “FOREIGN”.

```
table(z[, ncol(z1)])
```

```
##
##  0  1
## 82 18
```

```
table(z1[, ncol(z1)])/nrow(z1)
```

```
##
## DOMESTIC FOREIGN
##      0.82      0.18
```

Print out a summary of the whole dataframe.

```
summary(z1)
```

```
##           X1           X2           X3
## Min.      : 3.6518263   Min.      :-9.9976087   Min.      : 1.00
## 1st Qu.:12.4279359     1st Qu.: -5.4489063   1st Qu.: 5.00
## Median :16.8917099     Median : -2.1428592   Median : 6.00
## Mean     :16.9832838     Mean     : -1.1559632   Mean      : 6.19
## 3rd Qu.:21.2454707     3rd Qu.: 3.8595016   3rd Qu.: 7.25
## Max.     :34.3916371     Max.      : 9.6928934   Max.      :14.00
##           X4           X5           X6
## Min.      :0.0022537157   Min.      :0.00   Length:100
## 1st Qu.:0.0360274088     1st Qu.:1.00   Class :character
## Median :0.0723064742     Median :2.00   Mode  :character
## Mean     :0.1058796042     Mean      :2.33
## 3rd Qu.:0.1570614776     3rd Qu.:3.00
## Max.     :0.5629222776     Max.      :6.00
```

- Let $n = 50$. Create a $n \times n$ matrix R of exactly 50% entries 0's, 25% 1's 25% 2's. These values should be in random locations.

```

n <- 50

z <- c(rep(0, n * n / 2), rep(1, n * n / 4), rep(2, n * n / 4))

```



```

z <- sample(z)
R <- matrix(z,nrow = n, ncol = n)

head(R)

```

```

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
## [1,]    0    0    0    1    0    2    0    0    2    0    2    0    1    2
## [2,]    1    0    1    0    0    1    1    0    0    0    2    1    0    2
## [3,]    0    0    2    0    0    2    1    1    1    2    1    1    1    1
## [4,]    0    1    1    0    1    0    0    0    2    0    0    1    0    0
## [5,]    2    0    0    0    0    0    0    1    0    2    0    2    1    0
## [6,]    0    1    2    2    0    1    0    2    2    0    0    0    0    1
##      [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
## [1,]      2     0     2     1     1     0     0     0     0     1     0     1
## [2,]      1     1     0     2     0     0     2     1     0     1     2     2
## [3,]      2     2     1     0     2     1     0     2     1     2     1     2
## [4,]      1     0     1     1     0     0     0     0     0     2     1     0
## [5,]      0     0     1     1     1     0     1     1     2     2     0     0
## [6,]      1     0     1     0     0     2     0     2     2     1     0     0
##      [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
## [1,]      2     0     1     0     1     0     0     0     0     1     1     1
## [2,]      2     1     0     2     2     0     0     2     0     0     0     0
## [3,]      1     1     0     0     0     0     1     2     0     2     0     1
## [4,]      0     0     1     2     0     1     0     2     2     2     1     1
## [5,]      0     2     2     0     0     1     2     1     2     2     0     2
## [6,]      1     0     0     0     2     2     2     0     0     1     1     0
##      [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
## [1,]      1     0     2     0     2     0     1     0     2     0     0     0
## [2,]      2     1     1     2     0     2     0     0     2     0     0     0
## [3,]      1     0     0     2     1     0     1     0     0     1     0     1
## [4,]      0     0     0     2     0     0     0     0     0     1     2     0
## [5,]      0     0     2     0     0     2     0     0     0     1     0     0
## [6,]      2     0     1     1     0     0     1     0     2     1     0     2

```

- Randomly punch holes (i.e. NA) values in this matrix so that an each entry is missing with probability 30%.

```

for (i in 1:n) {
  for (j in 1:n) {
    if (runif(1) <= .30) {
      R[i,j] <- NA
    }
  }
}

head(R)

```

```

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
## [1,]   NA   NA    0    1   NA    2    0    0   NA   NA   NA    0    1   NA
## [2,]    1    0    1    0   NA   NA    1    0    0    0    2   NA    0    2
## [3,]    0    0    2    0    0    2    1   NA    1    2   NA    1    1   NA
## [4,]   NA    1   NA   NA    1    0    0    0    2    0    0   NA    0   NA

```

```
## [5,] NA 0 0 NA 0 0 0 1 0 2 0 2 1 NA
## [6,] 0 1 2 2 0 NA 0 2 NA 0 0 0 NA NA
##      [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
## [1,] 2 0 NA NA 1 0 0 0 NA NA 0 1
## [2,] 1 1 0 2 0 0 2 1 0 1 2 2
## [3,] 2 2 1 0 2 NA 0 2 1 2 NA 2
## [4,] 1 NA 1 1 0 0 NA 0 0 2 1 0
## [5,] 0 NA NA 1 1 0 1 1 2 NA NA NA
## [6,] 1 NA 1 0 0 NA 0 2 NA 1 0 0
##      [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
## [1,] NA NA 1 0 1 0 0 NA 0 1 1 1
## [2,] NA 1 0 2 NA 0 0 2 NA 0 NA 0
## [3,] NA NA 0 NA 0 NA 1 2 0 NA 0 1
## [4,] NA 0 NA NA 0 1 NA 2 NA NA NA 1
## [5,] NA NA 2 0 0 1 2 1 NA 2 NA NA
## [6,] 1 0 0 0 2 2 2 NA 0 NA NA 0
##      [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
## [1,] 1 0 NA NA NA 0 1 NA 2 0 NA 0
## [2,] 2 1 1 2 0 NA 0 NA NA NA NA 0
## [3,] NA NA NA 2 1 NA NA NA 0 1 NA 1
## [4,] 0 NA 0 2 0 NA 0 0 0 NA 2 0
## [5,] 0 NA 2 NA NA 2 NA 0 0 1 0 0
## [6,] 2 NA 1 1 0 0 1 NA 2 NA 0 NA
```

- Sort the rows in matrix R by the largest row sum to lowest. Be careful about the NA's!

```
z <- order(rowSums(R, na.rm = TRUE), decreasing = TRUE)
z
```

```
## [1] 23 26 3 30 39 34 36 37 8 2 28 45 13 15 21 24 48 14 20 31 9 17 6 7 10
## [26] 40 46 5 25 41 47 49 19 29 50 44 12 33 38 42 22 32 35 11 16 27 4 18 43 1
```

- We will now learn the `apply` function. This is a handy function that saves writing for loops which should be eschewed in R. Use the `apply` function to compute a vector whose entries are the standard deviation of each row. Use the `apply` function to compute a vector whose entries are the standard deviation of each column. Be careful about the NA's! This should be one line.

```
rows = apply(R, 1, sd, na.rm=TRUE)
rows2 = apply(R, 2, sd, na.rm=TRUE)

head(rows)
```

```
## [1] 0.67521402421 0.84334901040 0.82915619759 0.75377836144 0.83029750053
## [6] 0.84890218555
```

```
head(rows2)
```

```
## [1] 0.80229555709 0.65222452070 0.96430547933 0.90551882884 0.67680463660
## [6] 0.83333333333
```

- Use the `apply` function to compute a vector whose entries are the count of entries that are 1 or 2 in each column. This should be one line.

```
apply(R > 0, 2, sum, na.rm=TRUE)
```

```
## [1] 14 13 16 14 13 21 17 17 18 22 17 13 17 14 25 20 14 23 27 19 14 18 21 18 16
## [26] 16 12 10 17 19 17 13 17 23 14 15 22 17 16 22 22 24 15 17 20 15 18 14 16 13
```

- Use the `split` function to create a list whose keys are the column number and values are the vector of the columns. Look at the last example in the documentation `?split`.

```
a <- split(R, col(R))
head(a)
```

```
## $'1'
## [1] NA 1 0 NA NA 0 1 0 NA 0 2 NA NA 0 NA 0 0 NA 1 2 NA 0 1 2 NA
## [26] NA 0 1 0 2 0 1 0 0 NA 0 NA 1 2 NA NA 1 2 NA NA 0 NA NA 0 NA
##
## $'2'
## [1] NA 0 0 1 0 1 0 0 0 0 NA 0 0 NA NA 1 0 1 NA 2 0 NA 1 NA 1
## [26] 2 0 0 0 1 NA 0 0 NA NA 0 NA NA 0 NA 1 0 0 1 2 NA 0 0 1 0
##
## $'3'
## [1] 0 1 2 NA 0 2 2 0 NA NA 2 0 NA NA 2 NA NA NA 0 2 NA NA 2 0 1
## [26] 2 0 2 NA 0 NA 2 NA 0 NA NA NA 0 NA NA 1 NA NA 2 2 0 2 0 0 0
##
## $'4'
## [1] 1 0 0 NA NA 2 0 0 0 NA 0 0 2 2 0 NA 0 0 0 2 2 NA 0 NA 2
## [26] 0 1 0 1 0 0 NA 2 2 0 NA NA NA NA 0 2 NA 0 NA 1 NA 2 NA 0 NA
##
## $'5'
## [1] NA NA 0 1 0 0 NA 1 NA 0 2 1 1 NA 0 1 NA NA 1 1 NA 1 2 NA NA
## [26] 0 NA NA 0 0 2 0 NA 0 0 NA NA NA 1 0 NA 1 0 NA 0 0 0 0 NA 0
##
## $'6'
## [1] 2 NA 2 0 0 NA NA 0 NA NA 0 1 NA 0 1 1 0 0 1 NA NA 1 2 0 2
## [26] 0 2 1 2 NA 0 2 NA 0 0 1 2 1 NA 1 NA 2 0 0 1 2 NA 0 1 NA
```

- In one statement, use the `lapply` function to create a list whose keys are the column number and values are themselves a list with keys: “min” whose value is the minimum of the column, “max” whose value is the maximum of the column, “pct_missing” is the proportion of missingness in the column and “first_NA” whose value is the row number of the first time the NA appears.

```
func <- function(a) {
# as.list( #concatenation is a list
  c(min = min(a, na.rm = TRUE),
    max = max(a, na.rm = TRUE),
    pct_missing = mean(is.na(a)),
    first_NA = which.min(is.na(a))
  )
}

a <- lapply(split(R, col(R)), func)
head(a)
```

```
## $'1'
##      min      max pct_missing  first_NA
##      0.0      2.0          0.4        2.0
##
## $'2'
##      min      max pct_missing  first_NA
##      0.00     2.00          0.28        2.00
##
## $'3'
##      min      max pct_missing  first_NA
##      0.0      2.0          0.4         1.0
##
## $'4'
##      min      max pct_missing  first_NA
##      0.00     2.00          0.32        1.00
##
## $'5'
##      min      max pct_missing  first_NA
##      0.00     2.00          0.38        3.00
##
## $'6'
##      min      max pct_missing  first_NA
##      0.00     2.00          0.28        1.00
```

- Set a seed and then create a vector `v` consisting of a sample of 1,000 iid normal realizations with mean -10 and variance 100.

```
set.seed(1)
v <- rnorm(1000, -10, sqrt(100))
head(v)
```

```
## [1] -16.2645381074 -8.1635667578 -18.3562861241  5.9528080214 -6.7049222818
## [6] -18.2046838412
```

- Repeat this exercise by resetting the seed to ensure you obtain the same results.

```
set.seed(1)
v <- rnorm(1000, -10, sqrt(100))
head(v)
```

```
## [1] -16.2645381074 -8.1635667578 -18.3562861241  5.9528080214 -6.7049222818
## [6] -18.2046838412
```

- Find the average of `v` and the standard error of `v`.

```
mean(v)
```

```
## [1] -10.116481419
```

```
sd(v)/sqrt(1000)
```

```
## [1] 0.32726912404
```

- Find the 5%ile of v and use the `qnorm` function to compute what it theoretically should be. Is the estimate about what is expected by theory?

```
quantile(v, 0.05)
```

```
##           5%  
## -27.269599864
```

```
qnorm(.05, mean(v), sd(v))
```

```
## [1] -27.139332146
```

- What is the percentile of v that corresponds to the value 0? What should it be theoretically? Is the estimate about what is expected by theory?

```
c <- ecdf(v)(0)  
c
```

```
## [1] 0.842
```

```
#Generated  
d <- pnorm(0, mean(v), sd(v))  
d
```

```
## [1] 0.83584344242
```

```
#Theoretical  
d <- pnorm(0, -10, sqrt(100))  
d
```

```
## [1] 0.84134474607
```

#All three values generated are similarly. Therefore yes it is expected because v was generated using a